

Minimum Cut and Minimum k -Cut in Hypergraphs via Branching Contractions

Kyle Fox

joint with

Debmalya Panigrahi and Fred Zhang
Duke University

Talk Outline

- A (short) lecture, some recent work, and a new result...

Talk Outline

- A (short) lecture, some recent work, and a new result...
 1. Minimum cuts in graphs via branching random contractions [Karger, Stein '96]

Talk Outline


- A (short) lecture, some recent work, and a new result...
 1. Minimum cuts in graphs via branching random contractions [Karger, Stein '96]
 2. Minimum cuts in hypergraphs via *serial* random contractions
[Ghaffari *et al.* '17; Chandrasekharan *et al.* '18]

Talk Outline

- A (short) lecture, some recent work, and a new result...
 1. Minimum cuts in graphs via branching random contractions [Karger, Stein '96]
 2. Minimum cuts in hypergraphs via *serial* random contractions
[Ghaffari *et al.* '17; Chandrasekharan *et al.* '18]
 3. Faster minimum cuts in hypergraphs via *branching* contractions [FPZ]


Minimum Cuts

- A *cut* is a set of edges whose removal creates at least 2 connected components; a *k-cut* creates at least k components




Minimum Cuts

- A *cut* is a set of edges whose removal creates at least 2 connected components; a *k-cut* creates at least k components
- A *minimum (k-)cut* has as few edges as possible




Minimum Cuts

- A **cut** is a set of edges whose removal creates at least 2 connected components; a **k -cut** creates at least k components
- A **minimum (k -)cut** has as few edges as possible



Minimum Cuts

- A **cut** is a set of edges whose removal creates at least 2 connected components; a **k -cut** creates at least k components
- A **minimum (k -)cut** has as few edges as possible



Strategies

- We'll focus on finding minimum (2-)cuts (*i.e.*, minimum cuts)

Strategies

- We'll focus on finding minimum (2-)cuts (*i.e.*, minimum cuts)
- Given a graph with n vertices and m edges...


Strategies

- We'll focus on finding minimum (2-)cuts (*i.e.*, minimum cuts)
- Given a graph with n vertices and m edges...
 - Compute $n - 1$ minimum s,t -cuts in $O(mn^2)$ time

Strategies


- We'll focus on finding minimum (2-)cuts (*i.e.*, minimum cuts)
- Given a graph with n vertices and m edges...
 - ▶ Compute $n - 1$ minimum s,t -cuts in $O(mn^2)$ time
 - ▶ Guess an edge outside a minimum cut (uniformly at random) and **contract**...

Edge Contraction



- Identify endpoints and remove loops

Edge Contraction



- Identify endpoints and remove loops
- Like committing to *not* separating the endpoints with a cut

Contractions Preserve Cuts (Probably)

- **Lemma:** For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$

Contractions Preserve Cuts (Probably)

- **Lemma:** For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$
 - ▶ Each vertex is incident to at least $|C|$ edges


Contractions Preserve Cuts (Probably)

- **Lemma:** For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$
 - ▶ Each vertex is incident to at least $|C|$ edges
 - ▶ So there are at least $|C| (n / 2)$ edges

Repeating Contractions

[Karger '93]


- Repeatedly contract until two vertices remain



Repeating Contractions

[Karger '93]


- Repeatedly contract until two vertices remain



Repeating Contractions

[Karger '93]


- Repeatedly contract until two vertices remain



Repeating Contractions

[Karger '93]


- Repeatedly contract until two vertices remain



Repeating Contractions

[Karger '93]

- Repeatedly contract until two vertices remain



Repeating Contractions

[Karger '93]

- Repeatedly contract until two vertices remain
- Surviving edges form a cut



It Works!

- For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$

It Works!

- For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$

induction...

It Works!

- For *any* minimum cut C , we contract an edge of C with probability $\leq 2 / n$

induction...

- Minimum cut C survives the entire process with probability $\Omega(1 / n^2)$


Try, Try Again

- One run of the repeated contraction process finds a minimum cut with probability $\Omega(1 / n^2)$
- Run it $O(n^2)$ times to succeed with constant probability
- Or $O(n^2 \log n)$ times to succeed with *high probability* (probability $\geq 1 - 1 / n$)
- Takes $O(n^2)$ time to fully contract once, so $O(n^4 \log n)$ time total

Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain



Branching Contractions

[Karger, Stein '96]



- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain
- Make two copies of contracted graph, recurse, and return the smaller cut



Branching Contractions

[Karger, Stein '96]


- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain
- Make two copies of contracted graph, recurse, and return the smaller cut



Branching Contractions

[Karger, Stein '96]



- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain
- Make two copies of contracted graph, recurse, and return the smaller cut



Branching Contractions

[Karger, Stein '96]

- Probably for one run to succeed is too low!
- To boost probability, contract until $\lceil n/\sqrt{2} \rceil$ vertices remain
- Make two copies of contracted graph, recurse, and return the smaller cut



Still Fast

- A single try has the runtime recurrence $T(n) = 2T(n/\sqrt{2}) + O(n^2)$
- So $O(n^2 \log n)$ time per run

But More Accurate

- For *any* minimum cut C , probability we contract any edge in C before recursion is $\leq 1 / 2$

But More Accurate

- For *any* minimum cut C , probability we contract any edge in C before recursion is $\leq 1 / 2$

nasty induction...

But More Accurate

- For *any* minimum cut C , probability we contract any edge in C before recursion is $\leq 1 / 2$

nasty induction...

- We return a minimum cut with probability $\Omega(1 / \log n)$

Final Result for Graphs

- One run of the branching contraction process finds a minimum cut with probability $\Omega(1 / \log n)$

Final Result for Graphs

- One run of the branching contraction process finds a minimum cut with probability $\Omega(1 / \log n)$
- Run it $O(\log^2 n)$ times to succeed with high probability

Final Result for Graphs

- One run of the branching contraction process finds a minimum cut with probability $\Omega(1 / \log n)$
- Run it $O(\log^2 n)$ times to succeed with high probability
- $O(n^2 \log^3 n)$ time total!

Final Result for Graphs


- One run of the branching contraction process finds a minimum cut with probability $\Omega(1 / \log n)$
- Run it $O(\log^2 n)$ times to succeed with high probability
- $O(n^2 \log^3 n)$ time total!
- Faster than known deterministic algorithms when the graph is dense

Final Result for Graphs

- One run of the branching contraction process finds a minimum cut with probability $\Omega(1 / \log n)$
- Run it $O(\log^2 n)$ times to succeed with high probability
- $O(n^2 \log^3 n)$ time total!
- Faster than known deterministic algorithms when the graph is dense
- Can easily modify algorithm to find minimum k -cuts in $O(n^{2k-2} \log^2 n)$ time (best known algorithm)


Today's Real Topic: Hypergraphs

- A *hypergraph* is a collection of vertices and arbitrary subsets of vertices called *hyperedges*




Today's Real Topic: Hypergraphs

- A *hypergraph* is a collection of vertices and arbitrary subsets of vertices called *hyperedges*
- A *minimum (k-)cut* is still a minimum cardinality set of hyperedges whose removal creates at least k components




Today's Real Topic: Hypergraphs

- A *hypergraph* is a collection of vertices and arbitrary subsets of vertices called *hyperedges*
- A *minimum (k-)cut* is still a minimum cardinality set of hyperedges whose removal creates at least k components



Today's Real Topic: Hypergraphs

- A *hypergraph* is a collection of vertices and arbitrary subsets of vertices called *hyperedges*
- A *minimum (k-)cut* is still a minimum cardinality set of hyperedges whose removal creates at least k components



Hyperstrategies

- Again, focus on finding minimum (2-)cuts

Hyperstrategies


- Again, focus on finding minimum (2-)cuts
- Given a hypergraph with n vertices, m hyperedges, and total hyperedge size p

Hyperstrategies

- Again, focus on finding minimum (2-)cuts
- Given a hypergraph with n vertices, m hyperedges, and total hyperedge size p
 - Guess a hyperedge outside a minimum cut and *contract*...?


Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v




Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v
- Two copies of a path spanning over every vertex except v



Uniform Contracts are Bad

- $n - 1$ distinct edges of size $n - 1$, all containing vertex v
- Two copies of a path spanning over every vertex except v
- Minimum cut contains every hyperedge, but there is a $\geq 1/3$ probability of contracting a hyperedge each round



Really Bad

- Minimum cut survives all contractions with probability ***exponentially*** low in maximum hyperedge size [Kogan, Krauthgamer '15]
- Need to bias selection away from large hyperedges

Biased Selection

- Should contract each hyperedge e with probability proportional to $(1 - |e| / n)$ [Chandrasekaran, Xu, Yu '18]

Biased Selection

- Should contract each hyperedge e with probability proportional to $(1 - |e| / n)$ [Chandrasekaran, Xu, Yu '18]
- Minimum cut survives all contractions with probability $\Omega(1 / n^2)$!

Biased Selection

- Should contract each hyperedge e with probability proportional to $(1 - |e| / n)$ [Chandrasekaran, Xu, Yu '18]
- Minimum cut survives all contractions with probability $\Omega(1 / n^2)$!
- Can perform all contractions in $O(p \log n)$ time total [Ghaffari, Karger, Panigrahi '17]

Biased Selection

- Should contract each hyperedge e with probability proportional to $(1 - |e| / n)$ [Chandrasekaran, Xu, Yu '18]
- Minimum cut survives all contractions with probability $\Omega(1 / n^2)$!
- Can perform all contractions in $O(p \log n)$ time total [Ghaffari, Karger, Panigrahi '17]
- So compute a minimum cut with high probability in $O(p n^2 \log^2 n) = O(mn^3 \log^2 n)$ time

Biased Selection

- Should contract each hyperedge e with probability proportional to $(1 - |e| / n)$ [Chandrasekaran, Xu, Yu '18]
- Minimum cut survives all contractions with probability $\Omega(1 / n^2)$!
- Can perform all contractions in $O(p \log n)$ time total [Ghaffari, Karger, Panigrahi '17]
- So compute a minimum cut with high probability in $O(p n^2 \log^2 n) = O(mn^3 \log^2 n)$ time
- Minimum k -cut ($k \geq 3$) with high probability in $O(p n^{2k-1} \log n) = O(mn^{2k} \log n)$ time

Branching Contractions?

- Karger and Stein branch at carefully chosen graph orders

Branching Contractions?

- Karger and Stein branch at carefully chosen graph orders
- Branch too early, and the number of recursive subproblems becomes **too large**

Branching Contractions?


- Karger and Stein branch at carefully chosen graph orders
- Branch too early, and the number of recursive subproblems becomes **too large**
- Branch too late, and the probability of success becomes **too small**

Branching Contractions?

- Karger and Stein branch at carefully chosen graph orders
- Branch too early, and the number of recursive subproblems becomes **too large**
- Branch too late, and the probability of success becomes **too small**
- Large hyperedges make it impossible to branch at the same time they do


Many Branches Needed

- Same example as before, but replace the paths with two hyperedges containing every vertex except v




Many Branches Needed

- Same example as before, but replace the paths with two hyperedges containing every vertex except v




Many Branches Needed

- Same example as before, but replace the paths with two hyperedges containing every vertex except v
- Minimum cut contains every green hyperedge; contraction preserves cut with probability $\leq 2 / (n + 1)$



Many Branches Needed

- Same example as before, but replace the paths with two hyperedges containing every vertex except v
- Minimum cut contains every green hyperedge; contraction preserves cut with probability $\leq 2 / (n + 1)$
- Need $\Omega(n)$ independent contractions



A Smoother Procedure

- Larger hyperedges means we need more branches, but how many?

A Smoother Procedure

- Larger hyperedges means we need more branches, but how many?
- Math works out easiest if we could do a fractional number of branches before each contraction

A Smoother Procedure

- Larger hyperedges means we need more branches, but how many?
- Math works out easiest if we could do a fractional number of branches before each contraction
- Can achieve this ideal *in expectation* by branching with a probability dependent on the size of present hyperedges

Probabilistic Branching

[FPZ]

- Before *every* contraction, we flip coins to decide whether or not to copy the current hypergraph and recursively call our algorithm

Probabilistic Branching

[FPZ]

- Before *every* contraction, we flip coins to decide whether or not to copy the current hypergraph and recursively call our algorithm
- The larger the hyperedges, the more likely we are to branch

Bias Against Large Hyperedges

- Previous work biased hyperedge selection away from large hyperedges; meanwhile, we want large hyperedges to induce more branching

Bias Against Large Hyperedges

- Previous work biased hyperedge selection away from large hyperedges; meanwhile, we want large hyperedges to induce more branching
- Our algorithm moves all bias against large hyperedges into deciding whether or not to branch

Bias Against Large Hyperedges

- Previous work biased hyperedge selection away from large hyperedges; meanwhile, we want large hyperedges to induce more branching
- Our algorithm moves all bias against large hyperedges into deciding whether or not to branch
- First, we select a hyperedge uniformly at random and commit to contracting it

Bias Against Large Hyperedges

- Previous work biased hyperedge selection away from large hyperedges; meanwhile, we want large hyperedges to induce more branching
- Our algorithm moves all bias against large hyperedges into deciding whether or not to branch
- First, we select a hyperedge uniformly at random and commit to contracting it
- But before contraction, we copy the hypergraph with probability based on the selected hyperedge's size

Probability of Branching

- Intuitively, we want to balance the probability of contracting a minimum cut's hyperedge with the probability of branching

Probability of Branching

- Intuitively, we want to balance the probability of contracting a minimum cut's hyperedge with the probability of branching
- Then, *in expectation*, we preserve one copy of the minimum cut in either the contracted hypergraph or the copy we create before contraction

Probability of Branching

- **Lemma:** For *any* minimum cut C , an edge chosen uniformly at random is in C with probability $\leq \frac{1}{m} \sum_e \frac{|e|}{n}$

Probability of Branching

- **Lemma:** For *any* minimum cut C , an edge chosen uniformly at random is in C with probability $\leq \frac{1}{m} \sum_e \frac{|e|}{n}$
- So if we select hyperedge e , we branch with probability $|e| / n$

That Looks Familiar...

- We branch with probability $|e| / n$; Chandrasekaran *et al.* would have *selected* hyperedge e with probability proportional to $(1 - |e| / n)$

That Looks Familiar...

- We branch with probability $|e| / n$; Chandrasekaran *et al.* would have *selected* hyperedge e with probability proportional to $(1 - |e| / n)$
- Their algorithm is the same as selecting a hyperedge e uniformly at random, and then *redoing* the selection with probability $|e| / n$

The Algorithm

- Maintain a set S of hyperedges we believe belong to the a minimum cut in hypergraph H

The Algorithm

- Maintain a set S of hyperedges we believe belong to the a minimum cut in hypergraph H

BranchingContract(H, S):

Add each spanning hyperedge to S and remove it from H

If H has no edges, return S

Select hyperedge e uniformly at random

With probability $|e| / n$, return the smaller of the cuts


$\text{BranchingContract}(H / e, S)$ and

$\text{BranchingContract}(H, S)$

Otherwise, return $\text{BranchingContract}(H / e, S)$


Computation Tree

- Visualize the algorithm's execution as a rooted tree over hypergraphs; input hypergraph H is the root, each time we perform a contraction, that node gets a child




Computation Tree

- Our computation tree is probabilistic while Karger-Stein's is deterministic




our algorithm




Karger-Stein

Computation Tree

- Our computation tree is probabilistic while Karger-Stein's is deterministic
- They have more branches per level *in exception*



our algorithm



Karger-Stein

Just As Accurate

- We randomly select a hyperedge to contract and branch probabilistically

Just As Accurate

- We randomly select a hyperedge to contract and branch probabilistically

tedious induction...

Just As Accurate

- We randomly select a hyperedge to contract and branch probabilistically

tedious induction...

- We return a minimum cut with probability $\geq 1 / (2H_n - 2)$
where $H_n = 1 + 1/2 + \dots + 1/n = \Theta(\log n)$

And Nearly as Fast

- **Lemma:** Given a hypergraph of order n , the computation tree contains expected $O((n / n_0)^2)$ hypergraphs of order n_0

And Nearly as Fast

- **Lemma:** Given a hypergraph of order n , the computation tree contains expected $O((n / n_0)^2)$ hypergraphs of order n_0
- We give a method to contract any hyperedge e in $O(m(n - |e| + 1))$ time

And Nearly as Fast

- **Lemma:** Given a hypergraph of order n , the computation tree contains expected $O((n / n_0)^2)$ hypergraphs of order n_0
- We give a method to contract any hyperedge e in $O(m(n - |e| + 1))$ time
- Sum over all n_0 to get a total running time of $O(mn^2 \log n)$ for **BranchingContract**

And Nearly as Fast

- **Lemma:** Given a hypergraph of order n , the computation tree contains expected $O((n / n_0)^2)$ hypergraphs of order n_0
- We give a method to contract any hyperedge e in $O(m(n - |e| + 1))$ time
- Sum over all n_0 to get a total running time of $O(mn^2 \log n)$ for **BranchingContract**
- Can compute a minimum cut with high probability in $O(mn^2 \log^3 n)$ time, a nearly $\Omega(n)$ improvement

Low Rank

- *Rank* r of a hypergraph is the maximum hyperedge size

Low Rank

- **Rank** r of a hypergraph is the maximum hyperedge size
- We give a method to contract any hyperedge in $O(n^{r-1})$ time when r is a constant

Low Rank

- **Rank** r of a hypergraph is the maximum hyperedge size
- We give a method to contract any hyperedge in $O(n^{r-1})$ time when r is a constant
- Can compute a minimum cut with high probability in $O(n^r \log^2 n)$ time when $r \geq 3$

Low Rank

- **Rank** r of a hypergraph is the maximum hyperedge size
- We give a method to contract any hyperedge in $O(n^{r-1})$ time when r is a constant
- Can compute a minimum cut with high probability in $O(n^r \log^2 n)$ time when $r \geq 3$
- Nearly optimal for dense hypergraphs (and we match Karger-Stein when $r = 2$)

Minimum k -cut

- Branch more often when $k \geq 3$

Minimum k -cut

- Branch more often when $k \geq 3$
- Minimum k -cut with high probability in $O(mn^{2k-2} \log^2 n)$ time for any $k \geq 3$

Minimum k -cut

- Branch more often when $k \geq 3$
- Minimum k -cut with high probability in $O(mn^{2k-2} \log^2 n)$ time for any $k \geq 3$
- Minimum k -cut with high probability in $O(n^{2k-2} \log^3 n)$ time if $r = 2k - 2$ and $O(n^{\max\{r, 2k-2\}} \log^2 n)$ time otherwise for any constant $k, r \geq 2$

Thanks!